

Clay Pricing: Bring-Your-Own-Key vs. Clay-Managed Billing

Project Beacon — a detailed, factually sourced breakdown of how Clay charges for enrichment, and exactly what changes when you supply your own API keys instead of using Clay's managed keys

Compiled 2026-08-31 · Sources: clay.com/pricing, clay.com/faq (multiple), university.clay.com/docs (multiple), community.clay.com — all fetched live today. Builds on and corrects one factual claim in this project's prior *Clay_BrightData_Claude_Enrichment_Cost_Report_v2.pdf* (2026-08-24) — see §7.

What I could not do: you linked Clay's account-specific credit-usage settings page (`app.clay.com/workspaces/1342561/settings/credit-usage`). That page requires your logged-in session — I have no Clay credentials and no browser session for your workspace, and a direct fetch confirmed it just returns Clay's public marketing homepage, not your real usage data. Everything below is Clay's **published, public billing policy** — accurate and current, but it cannot tell you what *your* workspace has actually spent. §9 tells you exactly how to pull that number yourself in a few minutes.

1. Executive summary

6. BYOK requirements & what you give up

2. The two-currency system: Data Credits vs. Actions

7. Correction: how failed lookups are actually billed

3. Current plan tiers & pricing
8. Worked comparison: with your key vs. without

4. How AI/LLM enrichment columns are billed

9. Getting your workspace's real number

5. Bring-your-own-key: provider by provider

10. Recommendation

1. Executive summary

Clay bills on two separate meters — **Data Credits** (paying for data pulled from third-party providers) and **Actions** (paying for Clay's own platform work: running a column, calling an AI model, pushing a record out). Bringing your own API key for a given provider or AI model **eliminates the Data Credit charge for that provider entirely** — you pay that provider directly instead — but **Actions are still charged either way**, because Actions pay for Clay's orchestration, not the underlying data.

The practical effect: BYOK saves money specifically on **high-volume, expensive-per-unit lookups** (AI/LLM enrichment columns and premium data providers) *if* you already hold — or are willing to pay for — a sufficiently high-tier account with that provider directly. It does not make Clay's own platform fee (Actions) disappear, and for providers you'd only ever use a handful of credits on, the administrative overhead of managing a separate API key/billing relationship may not be worth it.

2. The two-currency system: Data Credits vs. Actions

CLAY'S OWN FAQ

Data Credits	Actions
--------------	---------

Pays for	Data purchased from Clay's marketplace of 150+ third-party providers (Apollo, Hunter, Findymail, Clearbit, People Data Labs, etc.) — includes AI/LLM enrichment credit cost	Clay's own platform orchestration: running a column, calling an AI model, pushing to a CRM, exporting a row
Baseline unit cost	Starts at \$0.05 each, decreases with volume	"Less than a cent each," cheaper at scale — no single published flat rate
Rolls over?	Yes — Launch/Growth plans can bank up to 2× the monthly allotment	No — resets every billing cycle
Bypassed by BYOK?	Yes — using your own key for that provider skips the Data Credit charge entirely	No — a successful BYOK enrichment still consumes 1 Action

Sources: clay.com/faq/what-are-actions-and-data-credits-how-do-they-work, university.clay.com/docs/actions-data-credits

3. Current plan tiers & pricing

CLAY'S OWN PRICING PAGE

Plan	Price	Data Credits/mo	Actions/mo
Free	\$0	100	500
Launch	\$167/mo annual · \$185/mo monthly	2,500–3,000*	15,000
Growth	\$446/mo annual · \$495/mo monthly	6,000	40,000
Enterprise	Custom (annual commitment)	100,000+	200,000+

***Unresolved discrepancy in Clay's own materials, flagged rather than silently picked:** clay.com/pricing states Launch includes **3,000** Data Credits; university.clay.com/docs/plans-and-billing states **2,500**. Both were fetched today. This is a ~17% swing on your single most-used tier — confirm the real number in your own account's Plan & Billing page (not the public docs) before budgeting against it.

Clay overhauled this entire pricing structure on 2026-03-11, cutting marketplace data costs 50–90% versus the prior model. Existing customers on legacy plans (Starter/Explorer/Pro) could keep those plans indefinitely, but the window to switch *between* legacy tiers closed 2026-04-10 — worth checking which regime your workspace is actually on, since a legacy plan's credit math won't match the table above at all.

Sources: clay.com/pricing, university.clay.com/docs/plans-and-billing, university.clay.com/docs/credit-usage

4. How AI/LLM enrichment columns are billed

CLAY'S OWN DOCS

This is the section most relevant to a Claude/OpenAI-heavy pipeline like this project's. Clay uses **two different pricing models** depending on which model a column calls:

Pricing type	Applies to	How it's charged
--------------	------------	------------------

Fixed	Most standard tasks, including Clay's own models (branded "Neon," "Helium," "Argon")	A flat number of Data Credits per task, regardless of actual token usage
Variable	Advanced reasoning / deep web-research models (the category Claude Opus/Sonnet with extended tool use falls into)	Data Credits based on actual token usage, plus a 20% premium over the underlying provider's own per-token price

Every AI enrichment — fixed or variable pricing — also consumes **1 Action** on top of whatever Data Credits it uses. Clay handles rate limits and API tier access automatically when you use its managed key.

Source: university.clay.com/docs/ai-tokens

5. Bring-your-own-key: provider by provider CLAY'S OWN DOCS

BYOK is a paid-tier feature (Launch plan or above), configured per-provider under Profile Settings → Connections or directly inside an enrichment column's settings. Confirmed supported providers:

Category	Providers
AI / LLM	OpenAI, Anthropic, Gemini
Company/contact data	Apollo
Email finding	Findymail, Prospeo
Email verification	Debounce, NeverBounce

For every provider on this list, connecting your own key means: **you skip that provider's Data Credit charge entirely** and are billed by that provider directly, at whatever your own plan/rate with them is.

Sources: university.clay.com/docs/clay-credit-conservation, university.clay.com/docs/use-ai-integration-overview

6. BYOK requirements & what you give up CLAY'S OWN DOCS

BYOK is not free of friction — three real costs come with it:

Provider	Minimum tier required
OpenAI ClayGent	Tier 2+ (≥450,000 tokens/min)
Anthropic ClayGent	Tier 4+ (≥400,000 tokens/min)
Gemini ClayGent	Tier 2+, Gemini 2.0 Flash only
Plain ChatGPT (non-ClayGent) columns	30,000 tokens/min, paid tier with GPT-4/GPT-4 Turbo access

These tier minimums are real gates, not soft suggestions. Anthropic Tier 4 in particular is not the default tier a new account gets — it's reached through Anthropic's own usage-based tier progression, which itself takes real spend and

time to unlock. If your organization's Anthropic account isn't already at Tier 4, BYOK for Anthropic-powered Clay columns isn't available to you today regardless of budget — this is worth checking on console.anthropic.com before assuming BYOK is an option.

Two further tradeoffs, both confirmed in Clay's own documentation:

1. **You lose Clay's built-in cost visibility.** "When using a personal API key, price breakdowns won't appear in the Clay UI. You'll need to monitor your usage and upgrade tiers manually" — with your own key, Clay stops showing you what that column is costing; you have to track it in the provider's own dashboard.
2. **Rate limiting becomes your problem.** With Clay's managed key, Clay handles rate limits and tier access automatically. With your own key, hitting your provider's rate limit is something you have to diagnose and fix yourself.

7. Correction: how failed lookups are actually billed

CORRECTS PRIOR REPORT

This project's prior report (§7, "SOURCED FROM CLAY DOCS") stated: "Failed lookups: Still billed. A 5-provider waterfall where 4 miss and the 5th hits still consumes 5 credits, not 1." Refetching Clay's own current FAQ today shows this is **no longer accurate** — and may not have been accurate even then; either Clay changed this policy in the March 2026 overhaul, or the prior report's source was misread. Do not carry the old claim forward.

Clay's FAQ ("Do unsuccessful searches consume Data Credits? What about Actions?"), quoted verbatim below because this exact wording matters:

"You'll only be charged Data Credits when Clay successfully returns data. The same is true of Actions – an enrichment only counts toward your overall platform usage when the platform actually delivers new information." ... "In waterfall sequences, each step of the waterfall that returns data will consume Data Credits (unless using API keys) and an Action." ... "if the entire waterfall returns no results, no Data Credits or Actions will be consumed."

Read precisely, this says: a **miss** on any individual provider inside a waterfall costs nothing; only the step that **actually returns data** is billed (one charge, from whichever provider in the chain finally succeeds) — not one charge per provider attempted along the way. If every provider in the waterfall misses, the entire lookup is free.

One place independent sources disagree with Clay's own FAQ wording: a Clay community post and a third-party pricing breakdown both describe cases where a provider attempt is charged then *refunded* only if it individually fails, implying a charge-then-reverse mechanic rather than never-charged-on-miss. Both descriptions may be reconciled (a refund and a no-charge look identical on your final bill) — but if your workspace's usage log ever shows a debit immediately followed by a matching credit for the same row, that's this mechanic, not a Clay error.

Source: clay.com/faq/do-unsuccessful-searches-consume-data-credits-what-about-actions (fetched verbatim, 2026-08-31)

8. Worked comparison: with your own key vs. without

Clay does not publish one universal "\$/lead" figure — cost depends entirely on which providers a given column's waterfall touches and how many steps it takes to get a hit. What follows is not a measurement of this project's actual usage (that requires §9); it applies the policy above to the concrete case an AI-heavy LinkedIn enrichment column represents.

Scenario	Data Credit cost	Action cost	Who's actually billing you
Clay-managed key — AI enrichment column, standard/fixed-price model	Flat Data Credit amount per successful task (rate not publicly listed per-model)	1 Action	Clay only
Clay-managed key — AI enrichment column, advanced/reasoning model (Claude Opus-class with extended tool use)	Actual token cost + 20% premium, converted to Data Credits	1 Action	Clay only

Your own key (BYOK) — either model type	\$0 to Clay	1 Action	Clay (Action fee) + Anthropic/OpenAI directly, at their list per-token price, no 20% premium
--	-------------	-------------	---

What this actually means for a Claude-heavy pipeline like this one: the 20% premium Clay adds to variable-priced/advanced-model calls is the single clearest, most quantifiable BYOK saving in this whole policy — it applies precisely to the "Claude Opus 5 + web_search" style column this project already uses and has already measured real token costs for (\$0.58/lead in the prior report, §8). If that column runs under Clay's managed key, the true Clay-side cost for the AI step alone is **that same real token spend × 1.20**, plus 1 Action, still payable to Clay; under BYOK it's the same token spend at Anthropic's own rate (no 20% markup) plus that same 1 Action to Clay. The Action fee itself (under a cent) is the same in both cases and is not what BYOK saves you — the 20% markup on tokens is.

What this does not tell you: whether Clay's *total* per-lead cost (AI columns + every data-provider column in the same waterfall — email, phone, company data, etc.) beats or loses to running Claude + web_search standalone, because that requires knowing exactly which non-AI columns ran on the same rows and at what Data Credit rate each charged. That number lives only in your workspace's usage panel — see §9.

9. Getting your workspace's real number

This is unchanged from the prior report's §9 — repeated here because it's still the only path to a real, defensible cost comparison, and it's what the credit-usage page you linked is actually for:

1. Open `app.clay.com/workspaces/1342561/settings/credit-usage` yourself (logged in) and read the Data Credit and Action totals for the billing period the POC ran in.
2. Open the exact table used for the POC and list every enrichment column that ran — Experience, Education, Languages, Courses, Projects, Summary, Headline each may be a separate column with its own provider and rate.
3. If the usage panel only shows workspace-wide totals rather than per-table, take a usage snapshot immediately before and after re-running the same columns on a small fresh batch (5–10 leads) to isolate this table's real consumption.
4. Divide by leads actually processed to get a real credits/lead and actions/lead figure — repeat on the ~20-lead POC batch, not just 1–2 leads, since per-lead cost varies with how many waterfall steps a given profile needs before something hits.
5. Multiply by your plan's real \$/credit (from your own Plan & Billing page, not the public docs — see the discrepancy flagged in §3) to get an actual \$/lead figure, directly comparable to the \$0.58/lead (Opus 5) and \$0.28/lead (Sonnet 5) Claude + web_search figures already measured in the prior report.
6. If any column already runs on a connected API key, separate those rows from the Data-Credit-billed rows before averaging — mixing the two produces a blended number that doesn't represent either billing model accurately.

10. Recommendation

Confirmed today, directly from Clay's own current documentation: Actions are billed regardless of whether you BYOK; only Data Credits are bypassed. BYOK's clearest, quantifiable win for this project specifically is avoiding the 20%

token-price premium on advanced-reasoning AI columns — everything else (data-provider BYOK, standard-model AI columns) depends on rates Clay doesn't publish per-provider.

Not yet confirmed: Anthropic BYOK requires your organization's Anthropic console account to already be at Tier 4 ($\geq 400,000$ TPM) — check console.anthropic.com before treating BYOK as available; it may not be, independent of any Clay-side decision. OpenAI BYOK similarly requires Tier 2+.

Corrected: do not reuse this project's prior claim that a waterfall bills every attempted provider on a miss — Clay's current FAQ says only the step that actually returns data is charged, and a total-miss waterfall is free.

Still open: your workspace's actual credit-usage numbers, which only exist behind your own login. \$9 is the fastest path to closing that gap — it takes minutes, not another research pass.

All sources fetched live on 2026-08-31: clay.com/pricing · clay.com/faq/what-are-actions-and-data-credits-how-do-they-work · clay.com/faq/what-are-clay-credits-and-how-do-they-work · clay.com/faq/do-unsuccessful-searches-consume-data-credits-what-about-actions · university.clay.com/docs/plans-and-billing · university.clay.com/docs/actions-data-credits · university.clay.com/docs/ai-tokens · university.clay.com/docs/clay-credit-conservation · university.clay.com/docs/credit-usage · community.clay.com (waterfall credits discussion, referenced not directly quoted). This project's own prior report, [Clay_BrightData_Claude_Enrichment_Cost_Report_v2.pdf](#) (2026-08-24), for the Claude+web_search cost figures cited in §8.